

**A.****1. Are foreign keys allowed to have null values?**

- a) No, foreign keys should have always a concrete value
- b) No, because the value of a key should be unique for any instance of a relation
- c) Yes, because there are situations in which such kind of information is not available
- d) Yes, foreign keys have always null values.

**2. What is *logical data independence*?**

- a) Changes made in physical schema of a database will not affect the conceptual schema
- b) Changes made in user interface will not affect the conceptual schema
- c) Changes made in conceptual schema will not affect the data access module.
- d) Changes made in conceptual schema of a database will not affect the external schemas

**3. Which of the following is not a criteria to use databases for your software system?**

- a) Preserve data integrity
- b) Huge amount of structured data
- c) Low-level data access
- d) Persistence

**4. Which of the following statements about dense-sparse indexes is *false*?**

- a) A dense index must be clustered.
- b) A dense index has at least one data entry for every search key which appears in the indexed file.
- c) A sparse index is typically much smaller than a dense index.
- d) A sparse index contains an entry for each page of records in a data file.

**5. Which of the following definitions is not an alternative of data entries in index files?**

- a) `<search_key, data record>`

2. What is logical data independence?

- a) Changes made in physical schema of a database will not affect the conceptual schema
- b) Changes made in user interface will not affect the conceptual schema
- c) Changes made in conceptual schema will not affect the data access module
- d) Changes made in conceptual schema of a database will not affect the external schemas

3. Which of the following is not a criteria to use database for your software system?

- a) Preserve data integrity
- b) High amount of structured data
- c) Low-level data access
- d) Persistence

4. Which of the following statement about dense-page indexes is false?

- a) A dense index must be clustered
- b) A dense index has at least one data entry for every search key which appears in the indexed file
- c) A sparse index is typically much smaller than a dense index
- d) A sparse index contains an entry for each page of records in a data file

5. Which of the following definition is not an alternative of data entries in index files?

- a) <search key, data record>
- b) <search key, primary key>
- c) <search key, record ID>
- d) <search key, list of record IDs>

6. The correct hierarchy of terms from largest to smallest is:

- a) field, record, table, database
- b) record, table, database, field
- c) database, table, field, record
- d) database, table, record, field

7. The purpose of a table is to:

- a) Present a document that can be posted on the internet
- b) Present a document in a professional fashion so it can be printed
- c) Store data about an entity type
- d) Answer a question about the database

8. In SQL, if it is needed to display records from lowest to highest value, which one of the following commands will be used?

- a) Sort ascending
- b) Remove filter
- c) Sort descending
- d) Referential integrity



Q<sub>1</sub>: SELECT DISTINCT R.B FROM R LEFT OUTER JOIN S ON R.B = S.A  
 Q<sub>2</sub>: SELECT DISTINCT S.A FROM R RIGHT OUTER JOIN S ON R.C = S.D

- a) Q<sub>1</sub> and Q<sub>2</sub> have the same number of tuples
- b) Q<sub>1</sub> and Q<sub>2</sub> have different number of tuples
- c) Number of tuples in Q<sub>1</sub> cannot be computed
- d) Number of tuples in Q<sub>2</sub> cannot be computed

12. Consider the following relation and functional dependencies:

S (A, B, C, D)                      {A → BCD, C → AB}

Which is the list of all candidate keys?

- a) A
- b) A, C
- c) AC
- d) AB

13. Which of the following best characterizes a balanced tree?

- a) All leaves are on the same level
- b) Leaves are positioned on at most 2 levels
- c) For each node the difference between its sub-trees heights is 1, 0 or -1
- d) All paths from root to leaves have the same length

14. Who is responsible for unpinning a page in buffer pool?

- a) The buffer manager
- b) The requestor of that page
- c) The frame which contains the page
- d) The replacement policy

15. The following SQL queries refer to relations R(a; b) and S(b; c):

Q<sub>1</sub>: SELECT \* FROM R INNER JOIN S ON R.b = S.b;

Q<sub>2</sub>: SELECT \* FROM R LEFT JOIN S ON R.b = S.b;

- a) Q<sub>1</sub> and Q<sub>2</sub> produce the same answer.
- b) The answer to Q<sub>1</sub> is always contained in the answer to Q<sub>2</sub>.
- c) The answer to Q<sub>2</sub> is always contained in the answer to Q<sub>1</sub>.
- d) Q<sub>1</sub> and Q<sub>2</sub> produce different answers.

9. Which of the following is not an alternative of mapping inheritance relationship between 2 classes A and B in tables from relational model?

- a) Create table A and de-normalize all attributes of B
- b) Create tables A and B, with their own attributes
- c) Create table B and de-normalize all attributes of A
- d) Create tables A, B and a cross-table between them.

10. Which of the following decompositions of the schema  $\{A, B, C, D, E, F\}$  are dependency preserving decompositions under the given functional dependencies?

- a)  $\{A, B, C, D, F\}$  and  $\{B, C, D, E\}$ : given  $ABC \rightarrow DEF$ ,  $BC \rightarrow EF$  and  $EF \rightarrow B$
- b)  $\{A, B, C, D, F\}$  and  $\{B, C, D, E\}$ : given  $AB \rightarrow CDE$ ,  $DE \rightarrow CF$  and  $C \rightarrow E$
- c)  $\{A, B, C, D, F\}$  and  $\{B, C, E\}$ : given  $ABC \rightarrow DEF$ ,  $BC \rightarrow E$  and  $C \rightarrow B$
- d)  $\{A, B, C, D, F\}$  and  $\{B, E, F\}$ : given  $ABC \rightarrow DEF$ ,  $BC \rightarrow F$  and  $EF \rightarrow B$

11. Given two relations,

$R:$

A	B	C
1	2	3
1	2	4
3	3	5

$S:$

A	D
1	4
1	7
3	8

and the following queries:

$Q_1$ : SELECT DISTINCT R.B FROM R LEFT OUTER JOIN S ON R.B = S.A

$Q_2$ : SELECT DISTINCT S.A FROM R RIGHT OUTER JOIN S ON R.C = S.D

- a)  $Q_1$  and  $Q_2$  have the same number of tuples
- b)  $Q_1$  and  $Q_2$  have different number of tuples
- c) Number of tuples in  $Q_1$  cannot be computed
- d) Number of tuples in  $Q_2$  cannot be computed

12. Consider the following relation and functional dependencies:

$S(A, B, C, D)$

$\{A \rightarrow BCD, C \rightarrow AB\}$

Which is the list of all candidate keys?

- a) A
- b) A, C
- c) AC
- d) AB

13. Which of the following best characterizes a balanced tree?



16. In the following, you may assume relations R(a, b) and S(b, c) have no NULL's, but may have duplicates.

(1) SELECT R.a FROM R, S

WHERE R.b = S.b

(2) SELECT R.a FROM R

WHERE R.b IN

(SELECT S.b FROM S)

a) Q<sub>1</sub> and Q<sub>2</sub> produce the same answer

b) The answer to Q<sub>1</sub> is always contained in the answer to Q<sub>2</sub>

c) The answer to Q<sub>2</sub> is always contained in the answer to Q<sub>1</sub>

d) Q<sub>1</sub> and Q<sub>2</sub> produce different answers

17. In the following, the results of Q<sub>1</sub> and Q<sub>2</sub> should be taken to be the result of the final SELECT \* FROM R.

Assume that the schema of relation R is R(a, b).

(1) UPDATE R SET b = 20 WHERE a = 10;

SELECT \* FROM R;

(2) DELETE FROM R WHERE a = 10;

INSERT INTO R VALUES(10, 20);

SELECT \* FROM R;

a) Q<sub>1</sub> and Q<sub>2</sub> produce the same answer

b) The answer to Q<sub>1</sub> is always contained in the answer to Q<sub>2</sub>

c) The answer to Q<sub>2</sub> is always contained in the answer to Q<sub>1</sub>

d) Q<sub>1</sub> and Q<sub>2</sub> produce different answers

18. In the following, R has attribute A, but its schema is otherwise not specified, nor is R relevant

(1) SELECT COUNT(DISTINCT b) FROM R;

(2) SELECT COUNT(b) FROM R;

a) Q<sub>1</sub> and Q<sub>2</sub> produce the same answer

b) The answer to Q<sub>1</sub> is always contained in the answer to Q<sub>2</sub>

c) The answer to Q<sub>2</sub> is always contained in the answer to Q<sub>1</sub>

d) Q<sub>1</sub> and Q<sub>2</sub> produce different answers

19. In the following expressions of relational algebra, the relation R has schema R(a, b);

(1)  $\pi_a(R) \times \pi_b(R)$

(2)  $\pi_{ab}(R \bowtie_{b=a} \rho(S)(a \rightarrow b, b \rightarrow a, R))$

a) Q<sub>1</sub> and Q<sub>2</sub> produce the same answer

b) The answer to Q<sub>1</sub> is always contained in the answer to Q<sub>2</sub>

c) The answer to Q<sub>2</sub> is always contained in the answer to Q<sub>1</sub>

d) Q<sub>1</sub> and Q<sub>2</sub> produce different answers

20. In the following the relations are R(a, b) and S(c, d).

(1) SELECT a

FROM R

WHERE b = ANY(SELECT d FROM S WHERE c = 5);

(2) SELECT a

FROM R

WHERE b = ANY(SELECT d FROM S WHERE c = 5);

a) Q<sub>1</sub> and Q<sub>2</sub> produce the same answer

b) The answer to Q<sub>1</sub> is always contained in the answer to Q<sub>2</sub>

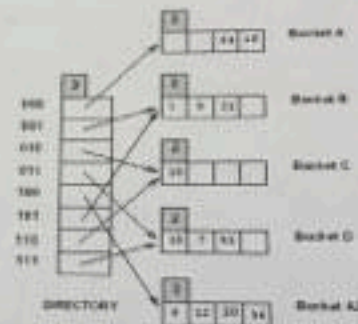
c) The answer to Q<sub>2</sub> is always contained in the answer to Q<sub>1</sub>

d) Q<sub>1</sub> and Q<sub>2</sub> produce different answers

A. (40%)

1	6	11	16
2	7	12	17
3	8	13	18
4	9	14	19
5	10	15	20

B. (15%) Consider the Extendible Hashing index shown below. Show the index after inserting with hash values 17 and 9.



C. (30%) Consider the following relation and its set of functional dependencies:

S(A, B, C, D, E, F, G)

{A → BC, ABC → D}

- List all of the candidate keys
- Give the minimal cover of the set of functional dependencies.
- Name the strongest normal form that is not violated by the relation containing attributes.
- Decompose R into a collection of BCNF relations if it is not in BCNF

D. (15%) Assume that you are given a relation schema R(A, B, C, D). Assume that r holds. Write an SQL query that checks whether the functional dependency

**C. (30%)** Consider the following relation and its set of functional dependencies:

$S(A, B, C, D, E, F, G)$

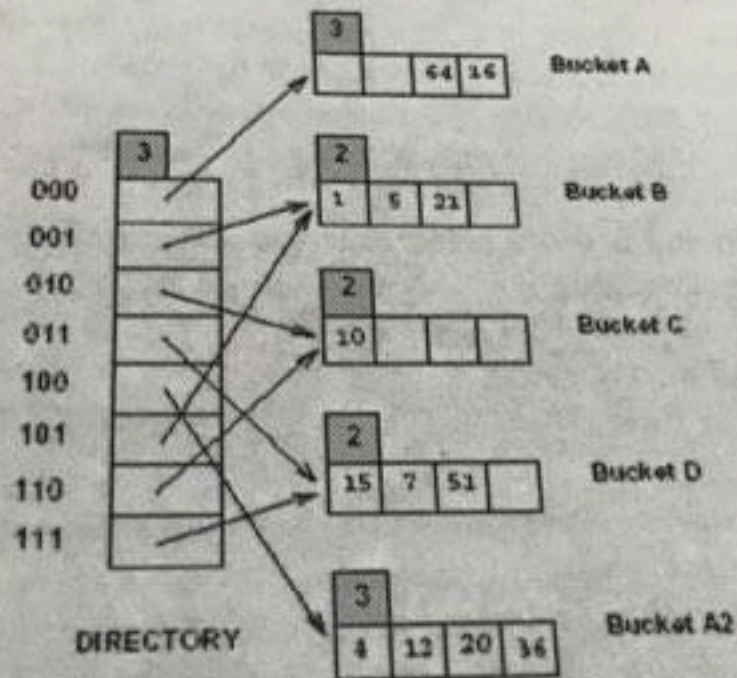
$\{A \rightarrow BC, ABC \rightarrow D\}$

- List all of the candidate keys.
- Give the minimal cover of the set of functional dependencies.
- Name the strongest normal form that is not violated by the relation containing these attributes.
- Decompose  $R$  into a collection of BCNF relations if it is not in BCNF

**D. (15%)** Assume that you are given a relation schema  $R(A, B, C, D)$ . Assume that no record has NULL values. Write an SQL query that checks whether the functional dependency  $AC \rightarrow BCD$  holds.



**B. (15%)** Consider the Extendible Hashing index shown below. Show the index after inserting entries with hash values 17 and 9.



**C. (15%)** Consider the following relation and its set of functional dependencies: